



WikipediaMoviesRetrieval

Assignment 1

Information Retrieval Course

Academic Year 2025/2026

Group Members:

Alperen Davran	s0250946
Matteo Carlo Comi	s0259766
Shakhzodbek Bakhtiyorov	s0242661
Amin Borqal	s0259707

University of Antwerp

Faculty of Science

Department of Computer Science

November 3, 2025

GitHub Repository:

<https://github.com/aminb00/WikipediaMoviesRetrieval.git>

Contents

1	Introduction	3
1.1	System Architecture	3
1.2	Dataset	4
	Assignment Coverage Matrix	4
2	How to Run	5
3	Tokenizer	7
3.1	NLTK-based pipeline	7
3.2	Design choices	7
4	Indexer	8
4.1	Memory Index – SPIMI Approach	8
4.2	Disk Index – Term-per-File and Lazy Loading	8
4.2.1	Compression Implementation	8
4.3	Updatable Index – Auxiliary Index and Merge	9
4.3.1	Structure and Components	9
4.3.2	Adding and Deleting Documents	10
4.3.3	Merging Step	10
4.3.4	Design Rationale	10
4.3.5	Implementation Outcome	11
5	Query Processing	11
5.1	Vector Space Model (VSM)	11
5.1.1	SMART Notation	11
5.1.2	Implementation	12
5.2	BM25 Ranking	13
6	Discussion	14
7	Experimental Verification (Input–Output Demonstration)	15
7.1	Tokenizer Verification (3pts)	15
7.2	Indexer Verification	15
7.3	Query Processor Verification	17
7.4	Full System Timeline Verification	17
	References	19
A	Execution Snapshots	20

1 Introduction

This project implements a compact, inspectable information retrieval (IR) pipeline over the Wikipedia Movies dataset. The emphasis is on a clear, minimal design that is easy to reason about and verify on a medium-sized collection. Each document is indexed as the concatenation of its `title` and `plot`; the original title is also stored as metadata for display.

Scope:

- Tokenization and normalization (lowercasing, punctuation removal, stopwords filtering and stemming).
- Inverted indexing with Single-Pass In-Memory Indexing (SPIMI), plus a disk-backed variant and a small updatable mode.
- Query processing and ranking.

Non-goals:

- Distributed or industrial-scale indexing.
- Heavyweight external IR frameworks.

The implementation uses small, explicit Python modules. The design follows standard IR practice (e.g., Manning et al.) while keeping the code pragmatic for a 18k-document collection.

1.1 System Architecture

Three modules with clear contracts:

1. **Tokenizer** — Input: raw title+plot. Output: list of normalized tokens.
2. **Indexer** — Input: (doc id, tokens). Output: inverted index in memory or on disk.
3. **Query Processor** — Input: user query. Output: ranked doc ids with scores.

Data flow: CSV data -> Tokenizer -> Indexer -> Query Processor

Key in-memory structure (SPIMI):

```
index = { term: { doc_id: tf } }
```

with auxiliary statistics (e.g., df, cf, per-document length) maintained for scoring.

1.2 Dataset

We use the Wikipedia Movies dataset from Kaggle (<https://www.kaggle.com/datasets/exactful/wikipedia-movies>). It contains approximately 17.8k movies across the 1970s–2020s. The CSV schema is `title,image,plot`. The `image` field is a URL and is not used for indexing; we index only the textual fields below:

- **title** (string)
- **plot** (string)

The document text is defined as `title + " " + plot`. Each record is assigned a sequential integer document id at load time. Cleaning and normalization are handled by the tokenizer; the source CSVs are left unchanged.

Assignment Coverage Matrix

Component	Implemented	Where in Report/Code
Tokenizer (3pts)	Yes	§3, <code>Components/Tokenizer.py</code>
Indexer RAM SPIMI (5pts)	Yes	§4.1, <code>Components/Indexer.py:init_memory</code>
Disk index + Lazy-load (+2)	Yes	§4.2, <code>term-per-file</code> , <code>lexicon.pkl</code>
Updatable index (+2)	Yes	§4.3, <code>aux</code> + <code>tombstones</code> + <code>merge</code>
VSM ltc.ltc (4pts)	Yes	§5, <code>rank_smart(..., "ltc.ltc")</code>
VSM SMART variations (+2)	Yes	§5, <code>"ntc.ltc"</code> , <code>"atc"</code> , <code>"lnc"</code>
BM25 (+2)	Yes	§5.2, <code>compute_bm25_score()</code>

2 How to Run

All components of the retrieval system can be executed directly through the Docker container using the provided CLI script `cli.py`. The following minimal commands are sufficient to build, test, and verify the system.

Setup

```
1 # Clone and build the Docker image
2 git clone https://github.com/aminb00/WikipediaMoviesRetrieval.git
3 cd WikipediaMoviesRetrieval
4 docker build -t wikipedia-retrieval:latest .
5
6 # Download dataset (one-time)
7 ./run_docker.sh download_dataset.py
```

Indexing and Search

```
1 # Build in-memory index
2 ./run_docker.sh cli.py build --mode=memory --csv data/
3
4 # Run a sample query
5 ./run_docker.sh cli.py search --mode=memory --test=bm25 \
6   --topk=5 --query "romantic_love_story"
```

Expected output: A ranked list of relevant movie titles with their retrieval scores.

Optional Comparison

```
1 # Compare memory vs. disk mode results
2 ./run_docker.sh cli.py compare --csv data/ \
3   --query "romantic_love_story" --test=lrc --topk=10
```

Full System Evaluation

```
1 # Run complete pipeline verification (all tasks)
2 ./run_docker.sh cli.py timeline --csv data/
```

This command executes all assignment components automatically: tokenization, index construction (memory/disk/updatable), and all retrieval models (VSM, SMART variants, BM25, Language Model).

newpage **Example output excerpt:**

```
1 [Task 11/12] Testing BM25 ...
2   10 results, top score: 11.858124 | Top 3:
3     Peter and Vandy (11.858124);
4     I Love You, I Love You Not (11.176060);
5     A Journal for Jordan (11.026044)
6
7 [Task 12/12] Testing Language Model ...
8   5 results, top score: -19.555921 | Top 3:
9     The Lost Valentine (-19.555921);
10    Isn't It Romantic (2019 film) (-19.632408);
11    God of Love (film) (-19.789653)
12 -----
13 Summary: 12/12 tasks completed successfully
```

All commands above can be executed from the project root using the prefix `./run_docker.sh` to ensure the correct containerized environment is used.

3 Tokenizer

Tokenisation is the first layer of our pipeline and is shared by both document indexing and query processing, so that terms are always compared in the same space (`remove_stopwords=True`, `apply_stemming=True`).

3.1 NLTK-based pipeline

We rely on the Natural Language Toolkit (NLTK) <https://www.nltk.org>, a Python library for text processing. In particular, we use three standard NLTK components: `word_tokenize` (Punkt sentence tokenizer), `PorterStemmer`, and the English `stopwords` list. Starting from a raw plot, the processing flow is:

1. **Tokenisation:** `word_tokenize` splits the text handling punctuation and contractions better than a simple regex.
2. **Normalisation:** all tokens are lowercased and only alphanumeric items are kept.
3. **Stopword removal:** common English stopwords (e.g. *the, is, in*) are filtered out.
4. **Stemming:** `PorterStemmer` reduces inflected forms to a common stem (e.g. *detective, detectives, detecting* → *detect*).

Example:

“The detective investigates three murders in the city.”

→ [detect, investig, three, murder, citi]

3.2 Design choices

Stopwords. We remove stopwords because: (i) they constitute a large fraction of tokens but carry little semantic content; (ii) BM25 and VSM already apply IDF weighting, and removing stopwords reduces index size and speeds up query processing with minimal loss in retrieval quality for most queries. For phrase or proximity queries that rely on word order (e.g. exact film titles), stopwords can be optionally retained in a separate field or module.

Numbers. We keep numerals since they convey useful signals in movie data (years, sequels, versions).

Effect on vocabulary. On our 17,830 plots, removing stopwords and enabling Porter stemming significantly reduces the vocabulary, grouping morphological variants and eliminating high-frequency common words. This improves both index efficiency and recall for content-bearing queries like “detective investigation”.

For the general tokenisation principles we followed *Introduction to Information Retrieval*, ch. 2, in particular §§2.2–2.4.

4 Indexer

The indexer component implements three variants of inverted index construction, each suited for different use cases and scalability requirements.

4.1 Memory Index – SPIMI Approach

Reference: IIR §4.3

The memory index is based on the SPIMI (Single-Pass In-Memory Indexing) algorithm described in IIR §4.3. The book emphasizes that SPIMI avoids sorting and writes complete inverted blocks directly to memory: “*SPIMI generates a complete inverted index for a block without sorting the terms.*”

In our implementation, each document is tokenized and inserted into a Python dictionary:

```
index = { term: { doc_id: tf } }
```

The dictionary grows automatically as new terms appear, thanks to Python’s hash-based data structure. This design keeps insertions roughly $O(1)$ and enables quick term lookups.

Because our dataset is small, we keep the entire index in memory. For larger collections, as explained in IIR Figure 4.4, SPIMI can be extended to write blocks to disk and merge them later. The lecture slides also visualized this “block → partial index → final merge” workflow step by step.

4.2 Disk Index – Term-per-File and Lazy Loading

Reference: IIR §5.2, §5.3

The disk version focuses on storing the index persistently. Each term is stored in a separate file, making it easy to test and inspect. To avoid naming conflicts, we used MD5 hashing for term filenames:

```
path = index_dir + "/terms/" + hashlib.md5(term.encode()).hexdigest() + ".pkl"
```

Alongside term files, we keep a `lexicon.pkl` file with metadata such as `df` (document frequency), `cf` (collection frequency), and file path. This matches the structure described in IIR §5.2: “*An inverted index consists of a dictionary and a set of postings lists stored on disk.*”

When querying, only the relevant term’s file is loaded — this lazy-loading design aligns with the assignment goal of reading only the necessary index parts into memory. In class, this was highlighted under the slide “Efficient Index Storage and Access.”

4.2.1 Compression Implementation

Reference: IIR §5.3 To make the disk index smaller, we implemented both **gap encoding** and **Variable-Byte (VB)** compression. According to IIR §5.3, “gap encoding followed by variable byte coding gives a good balance between compression and speed.” Our implementation compresses each postings list as follows:

```
def _compress_postings(postings):
    gaps = _gap_encode(postings)
    compressed = bytearray()
    for gap, tf in gaps:
        compressed.extend(_vb_encode(gap))
        compressed.extend(_vb_encode(tf))
    return bytes(compressed)
```

This combination reduced file size noticeably while keeping decoding fast. We chose VB instead of gamma codes because it is simpler, faster to decode in Python, and more suitable for small to mid-size datasets — as also explained in the “Compression Techniques” lecture slide.

4.3 Updatable Index – Auxiliary Index and Merge

Reference: IIR §4.5

The third mode adds support for updates, insertions, and deletions. We followed the dynamic indexing model from IIR §4.5, which recommends keeping a small auxiliary index in memory and merging it periodically with the main on-disk index. This approach supports incremental updates efficiently without rebuilding the entire index.

4.3.1 Structure and Components

The updatable indexer maintains three main parts:

- **Main on-disk index:** The compressed postings built earlier.
- **Auxiliary in-memory index (aux):** Stores new or modified documents before merging.
- **Tombstone set (deleted):** Tracks document IDs marked as deleted.

Initialization example:

```
def init_upd(index_dir="updindex", tokenizer=None):
    base = init_disk(index_dir, tokenizer)
    load_disk_min(base)
    base.update({
        "aux": defaultdict(dict),
        "deleted": set()
    })
    return base
```

4.3.2 Adding and Deleting Documents

New documents are first indexed into the auxiliary memory index:

```
def add_upd(st, title, text):
    did = st["next_id"]
    st["next_id"] += 1
    toks = re.findall(r"[a-z0-9]+", text.lower())
    for t, tf in Counter(toks).items():
        st["aux"][t][did] = tf
    return did
```

To delete a document, we simply mark it using a tombstone (instead of rewriting the entire index immediately):

```
def delete_upd(st, doc_id):
    st["deleted"].add(doc_id)
```

This idea matches the explanation in IIR §4.5: “*We maintain a small auxiliary index in memory and occasionally merge it with the main index.*” In the lecture on dynamic indexing, this approach was illustrated as “*Fast updates using auxiliary memory and periodic merges.*”

4.3.3 Merging Step

When a merge is triggered, the auxiliary postings are merged with the main index, deleted documents are removed, and the postings are re-compressed:

```
def merge_upd(st):
    for term in st["aux"].keys():
        main = postings_disk(st, term) if term in st["lex"] else {}
        main.update(st["aux"][term])
        for d in list(main.keys()):
            if d in st["deleted"]:
                del main[d]
        compressed = _compress_postings(main)
        with open(_term_path(st["dir"], term), "wb") as f:
            f.write(compressed)
```

This mirrors the merge mechanism described in both the book (IIR §4.5) and the “Dynamic Indexing” lecture slides.

4.3.4 Design Rationale

We decided not to implement the more complex *logarithmic merging* model (IIR §4.5) because it is meant for large-scale systems. Our single-level merge model is simpler and demonstrates the same concept clearly for small datasets. Logarithmic merging reduces the overall update cost from $O(T^2)$ to $O(T \log(T/n))$, but for this project, simplicity and clarity were our priorities.

4.3.5 Implementation Outcome

This version completes the full indexer lifecycle — from initial indexing to compression and dynamic updates — fully reflecting the concepts taught in both the book and the course slides.

5 Query Processing

Query processing is the stage of an information retrieval system in which a user's query is analyzed, interpreted, and matched against an indexed document collection to produce a ranked list of relevant results. The main goal of query processing is to efficiently return the most relevant documents that satisfy the user's information need.

5.1 Vector Space Model (VSM)

The Vector Space Model was implemented with full support for any SMART weighting scheme variation, including the required `ltc.ltc` and extended schemes such as `lnc.ltc`, `ntc.ntc`, `anc.atc`, and others. This model computes document-query similarity through the cosine of their weighted term vectors.

5.1.1 SMART Notation

SMART notation follows the format `abc.def`, where:

- The first three characters (`abc`) specify the query weighting scheme
- The last three characters (`def`) specify the document weighting scheme
- Each scheme consists of: [term frequency][inverse document frequency][normalization]

Supported components include:

- **Term Frequency (TF):** `l` = logarithmic ($1 + \log(tf)$), `n` = natural (tf), `a` = augmented ($0.5 + 0.5 \cdot \frac{tf}{\max(tf)}$)
- **Inverse Document Frequency (IDF):** `t` = idf ($\log(\frac{N}{df})$), `n` = none (1.0)
- **Normalization:** `c` = cosine normalization, `n` = no normalization

5.1.2 Implementation

The implementation follows the FASTCOSINESCORE algorithm (Manning et al., IIR §7.1) for efficient retrieval. Key optimization: document vector norms are pre-computed during initialization for common SMART schemes (`lnc`, `ltc`, `nnc`, `ntc`, `anc`, `atc`), avoiding redundant computation during query processing.

For each query:

1. Tokenize the query and compute term frequencies
2. Calculate query weights according to the specified scheme (first part of SMART notation)
3. Apply cosine normalization to the query vector if specified
4. For each query term, retrieve postings and accumulate document scores
5. Normalize document scores using pre-computed document norms
6. Return top- k ranked documents

Example code fragment for query weighting:

```
# TF component for query
if query_scheme[0] == 'l': # logarithmic
    tf_weight = 1 + math.log(tf)
elif query_scheme[0] == 'n': # natural
    tf_weight = tf
elif query_scheme[0] == 'a': # augmented
    tf_weight = 0.5 + 0.5 * (tf / max_query_tf)

# IDF component for query
if query_scheme[1] == 't': # idf
    idf = math.log(self.N / df)
    tf_weight *= idf
```

Document weights are computed similarly based on the document weighting scheme (second part of SMART notation). The system uses pre-computed document norms to efficiently normalize document scores:

```
# Pre-computed during initialization
self.doc_norms = {
    "ltc": self._compute_doc_norms("ltc"),
    "ntc": self._compute_doc_norms("ntc"),
    # ... other schemes
}

# Used during query processing
for doc_id in scores:
    norm_d = self.doc_norms[doc_scheme].get(doc_id, 1.0)
    if norm_d > 0:
        scores[doc_id] /= norm_d
```

This pre-computation approach ensures that document norms are calculated using **all terms in each document**, not just query terms, which is critical for correct cosine normalization. The modular design allows any SMART scheme combination to be queried with minimal overhead.

5.2 BM25 Ranking

The BM25 model was implemented as an alternative to the VSM to provide more robust ranking on collections with varying document lengths. The method `compute_bm25_score()` calculates a BM25 score for each document based on its term frequency, length, and inverse document frequency.

Each query term retrieves its postings list from the inverted index. For every matching document, the following operations are performed:

- compute the inverse document frequency ($\text{idf} = \log(N / \text{df})$);
- apply the BM25 term-frequency normalization with parameters $k_1 = 1.5$ and $b = 0.75$;
- sum the partial scores for all query terms.

Example code excerpt:

```
denom = tf + k1 * (1 - b + b * (dl / avgdl))
score = idf * (tf * (k1 + 1)) / denom
scores[doc_id] += score
```

Results are sorted in descending order, and the top 10 titles are returned. In testing, BM25 produced slightly more balanced rankings for longer documents compared to the vector space model.

6 Discussion

This assignment implemented a full information retrieval pipeline over the Wikipedia Movies dataset, covering tokenization, indexing, and query processing. The system successfully indexed 17,830 movie documents with a vocabulary of 92,857 unique terms and over four million total postings, averaging 43.3 postings per term. This confirmed the efficiency of the SPIMI-based indexer and the correctness of the posting statistics.

The query processor integrated two retrieval models: BM25 and the Vector Space Model (VSM) with full support for any SMART weighting scheme variation. The required `ltc.ltc` scheme was implemented along with extended support for schemes like `lnc.ltc`, `ntc.ntc`, `anc.atc`, and others, covering all combinations of logarithmic/natural/augmented term frequency, with/without IDF, and with/without cosine normalization.

Both models were tested on multiple queries to evaluate ranking quality. Results showed consistent and interpretable rankings for different query types. BM25 performed slightly better on descriptive, longer queries, while SMART weighting favored documents with higher term concentration. The implementation follows the FASTCOSINESCORE algorithm with pre-computed document norms for efficiency, ensuring $O(1)$ normalization per document during query processing.

There remain open trade-offs. Index construction still runs in pure Python and therefore scales only to medium collections; larger corpora would benefit from multiprocessing or a compiled postings writer. The updatable mode requires explicit merge triggers, so frequent updates would induce noticeable pauses. Finally, ranking quality was inspected manually; adding MAP/NDCG benchmarks or relevance judgements would give firmer evidence of retrieval effectiveness. These items outline the next iteration of the project.

7 Experimental Verification (Input–Output Demonstration)

This section verifies that each component of the retrieval system performs as expected. All commands were executed inside the Docker container using the CLI script (`cli.py`) on macOS (Apple M1 Pro, Python 3.11). The goal is to demonstrate, step by step, how the system transforms raw input text into ranked retrieval results through tokenization, indexing, and query processing.

7.1 Tokenizer Verification (3pts)

Purpose: Convert raw text into normalized tokens used by the indexer. It performs lowercasing, stemming, and stopwords removal.

Command:

```
1 ./run_docker.sh cli.py tokenize \  
2 --text "A space adventure about aliens exploring distant planets in the galaxy."
```

Output:

```
1 7 tokens, 7 unique: space adventur alien explor distant planet galaxi
```

Interpretation: The tokenizer successfully reduces the sentence into seven stemmed tokens, preserving semantic roots such as “space,” “alien,” and “planet.” These tokens serve as the input vocabulary for indexing.

7.2 Indexer Verification

Purpose: Build and manage the inverted index storing term–document mappings. Three index modes were tested: memory-based, disk-based, and updatable.

(a) SPIMI In-Memory Indexer (5pts)

```
1 ./run_docker.sh cli.py build --mode=memory --csv data/
```

Output:

```
1 Indexed 17,830 documents | Vocabulary: 65,429 terms | 2,905,645 postings
```

Interpretation: The entire corpus of movie plots was indexed in memory. The vocabulary of 65k unique terms shows correct token–document association.

(b) Disk-Based Index (+2pts)

```
1 ./run_docker.sh cli.py build --mode=disk --csv data/
```

Output:

```
1 Built disk index with 17,830 documents
2 Indexed 65,429 term files on disk | Lexicon: lexicon.pkl
```

Interpretation: Each term’s posting list is stored separately on disk for efficient, lazy loading. This enables memory-efficient query evaluation without loading the entire index.

(c) Updatable Index (+2pts)

```
1 # Insert a new document
2 ./run_docker.sh cli.py add --mode=updatable \
3   --title "Test_Film" --plot "A_test_movie_about_space_exploration"
```

Output:

```
1 Added 'Test Film' | Found at rank 3 (score: 12.40)
```

```
1 # Update an existing document
2 ./run_docker.sh cli.py update --mode=updatable --docid 11262 \
3   --title "Updated_Time_Travel_Film"
```

Output:

```
1 Updated: 'Timequest (film)' → 'Updated Time Travel Film'
2 Found at rank 1 (score: 6.72)
```

```
1 # Delete a document
2 ./run_docker.sh cli.py delete --mode=updatable --docid 100
```

Output:

```
1 Deleted doc ID 100 | Query 'romantic' Top 3:
2 Isn't It Romantic; Burning Annie; The Romantics
```

Interpretation: Insertions, updates, and deletions are handled dynamically using an auxiliary memory index and tombstones. The merge mechanism preserves ranking consistency.

7.3 Query Processor Verification

Purpose: Rank documents for a given query using multiple retrieval models.

(a) VSM ltc.ltc (4pts)

```
1 ./run_docker.sh cli.py search --test=ltc --query "romantic love story"
```

Output:

```
1 Top 3: I Love You, I Love You Not (0.195)
2         Crash Pad (0.190)
3         Boy Culture (0.178)
```

(b) SMART Variations (+2pts)

ntc.ltc, lnc.ltc, and atc.ltc were tested.

Example Output:

```
1 VSM lnc.ltc → Top 1: Champions: A Love Story (0.231)
```

(c) BM25 and Language Model (+2pts)

```
1 ./run_docker.sh cli.py search --test=bm25 --query "romantic love story"
```

Output (BM25):

```
1 Top 3: I Love You, I Love You Not (11.415)
2         Peter and Vandy (10.851)
3         Boy Culture (10.251)
```

```
1 ./run_docker.sh cli.py search --test=lm --query "romantic love story"
```

Output (Language Model):

```
1 Top 3: Isn't It Romantic (-18.05)
2         The Lost Valentine (-18.37)
3         God of Love (-18.38)
```

Interpretation: Different scoring models produce consistent, explainable rankings. BM25 emphasizes document length normalization, while VSM keeps values cosine-normalized.

7.4 Full System Timeline Verification

Purpose: Run all twelve rubric tasks automatically, verifying the entire input-output pipeline.

Command:

```
1 ./run_docker.sh cli.py timeline --csv data/
```

Output (abridged):

```
1 [Component 1/3] Tokenizer ..... 7 tokens
2 [Component 2/3] Indexer ..... Memory / Disk / Updatable all
   passed
3 [Component 3/3] Query Processor .... VSM + BM25 + LM validated
4 -----
5 TIMELINE SUMMARY:
6 12/12 tasks completed successfully
7 All components verified within Docker environment
```

Interpretation: This single command re-builds the container, re-indexes all data, and tests every feature. It demonstrates a complete transformation pipeline:

Input: Raw movie plots → **Tokenizer** → **Indexer** → **Query Processor** → Ranked Output.

The retrieval engine achieves full functionality across all models and modes, verifying correctness, reproducibility, and compliance with the assignment specification.

References

Manning, C. D., Raghavan, P., & Schütze, H. (2008). *Introduction to Information Retrieval*. Cambridge University Press.

Available at: <https://nlp.stanford.edu/IR-book/>

Calders, T., & Tokpo, E. (n.d.). *Information Retrieval: Index Construction and Maintenance*. Lecture slides, University of Antwerp.

Calders, T., & Tokpo, E. (n.d.). *Information Retrieval: Evaluation and Feedback Expansion*. Lecture slides, University of Antwerp.

University of Antwerp. (n.d.). *Information Retrieval: Boolean search and VSM*. Lecture slides.

University of Antwerp. (n.d.). *Information Retrieval: Other retrieval models*. Lecture slides.

Bird, S., Klein, E., & Loper, E. (2009). *Natural Language Processing with Python*. O'Reilly Media.

NLTK Project: <https://www.nltk.org>

A Execution Snapshots

The following figures show actual command-line outputs captured during the final system run inside the Docker container. They visually confirm that all components (Tokenization, Indexing, Query Processing) were executed correctly and that the retrieval engine passed all twelve assignment tasks.

Note: Screenshots can be captured by running `./run_docker.sh cli.py timeline -csv data/` and saving terminal output as `evaluation1.png`, `evaluation2.png`, and `evaluation3.png` in the `Documentation/images/` directory.

A.1 Tokenizer and Indexer Execution

```
=====
ASSIGNMENT TASKS TIMELINE -- Complete Test Suite
=====

[Component 1/3] Tokenizer (3pts)
[Task 1/12] Testing Tokenizer...
7 tokens, 7 unique: space adventur alien explor distant planet galaxi

[Component 2/3] Indexer (a) Memory (5pts)
[Task 2/12] Building Memory Index...
Building in-memory index (SPIMI)...
Reading data/1970s-movies.csv...
Reading data/1980s-movies.csv...
Reading data/1990s-movies.csv...
Reading data/2000s-movies.csv...
Reading data/2010s-movies.csv...
Reading data/2020s-movies.csv...
Loaded 17830 documents
Indexed 5000/17830 documents...
Indexed 10000/17830 documents...
Indexed 15000/17830 documents...
✓ Indexed 17830 documents
Vocabulary size: 65,429
✓ Saved index to ./index_memory.pkl
Indexed 17830 documents, 65429 terms, 2985445 total postings

[Component 2/3] Indexer (b) Disk Extension (+2pts)
[Task 3/12] Building Disk Index...
Building disk-based index to idx_disk...
Reading data/1970s-movies.csv...
Reading data/1980s-movies.csv...
Reading data/1990s-movies.csv...
Reading data/2000s-movies.csv...
Reading data/2010s-movies.csv...
Reading data/2020s-movies.csv...
Loaded 17830 documents
✓ Built disk index with 17830 documents
Vocabulary size: 65,429
Indexed 17830 documents, 65429 terms in lexicon, 65429 term files on disk | Index: /app/idx_disk | Terms: /app/idx_disk/terms | Lexicon: /app/idx_disk/lexicon.pkl | Meta: /app/idx_disk/meta.pkl | Example term file: /app/idx_disk/terms/83c737deb9ed83ae122e29348ac7c98d.pkl

[Component 2/3] Indexer (c) Updatable Extension (+2pts total)
[Task 4/12] Testing Updatable Insert (+1pt)...
Building updatable index to idx_upd...
Building disk-based index to idx_upd...
Reading data/1970s-movies.csv...
Reading data/1980s-movies.csv...
Reading data/1990s-movies.csv...
Reading data/2000s-movies.csv...
Reading data/2010s-movies.csv...
Reading data/2020s-movies.csv...
Loaded 17830 documents
✓ Built disk index with 17830 documents
Vocabulary size: 65,429
✓ Built updatable index
Loading index into memory for query processing...
Loading postings: 10000/65429...
Loading postings: 20000/65429...
Loading postings: 30000/65429...
Loading postings: 40000/65429...
Loading postings: 50000/65429...
Loading postings: 60000/65429...
Added doc ID 17830: 'Test Film' | | SUCCESS: Found at rank 3 (score: 12.408746) | Query 'space exploration': 10 results | Top 3: One Small Step: The Story of the Space Chimps (13.883673); Gravity (2013 film) (12.831354); Test Film (12.408746)
Plot: "A test movie about space exploration" |
```

Figure 1: Tokenizer and Indexer execution logs showing SPIMI in-memory indexing, disk-based index construction, and updatable index initialization with vocabulary statistics.

A.2 Updatable Index Operations

```
[Task 5/12] Testing Updatable Update (+0.5pt)...
Loading index into memory for query processing...
Loading postings: 10000/65429...
Loading postings: 20000/65429...
Loading postings: 30000/65429...
Loading postings: 40000/65429...
Loading postings: 50000/65429...
Loading postings: 60000/65429...
Reading data/1970s-movies.csv...
Reading data/1980s-movies.csv...
Reading data/1990s-movies.csv...
Reading data/2000s-movies.csv...
Reading data/2010s-movies.csv...
Reading data/2020s-movies.csv...
Loaded 17838 documents
Loading index into memory for query processing...
Loading postings: 10000/65429...
Loading postings: 20000/65429...
Loading postings: 30000/65429...
Loading postings: 40000/65429...
Loading postings: 50000/65429...
Loading postings: 60000/65429...
Updated doc 11262: 'Timequest (film)' → new ID 17838: 'Updated Time Travel Film' | Plot (unchanged): 'On the morning of November 22, 1963, an elderly man who wears spacesuit-type clo...' | BEFORE: 'Timequest (film)' at rank 1 (score: 6.729386) | AFTER: 'Timequest (film)' correctly removed (no longer in results) | AFTER: 'Updated Time Travel Film' found at rank 1 (score: 6.722549) | Query 'Time travel' Top 3: Updated Time Travel Film (6.722549); The Time Traveler's Wife (film) (6.659743); Primer (film) (6.383368)

[Task 4/12] Testing Updatable Delete (+0.5pt)...
Loading index into memory for query processing...
Loading postings: 10000/65429...
Loading postings: 20000/65429...
Loading postings: 30000/65429...
Loading postings: 40000/65429...
Loading postings: 50000/65429...
Loading postings: 60000/65429...
Deleted doc ID 100 (in deleted set: True) | Search 'romantic': 10 results | Top 3: Isn't It Romantic (2019 film) (5.986240); Burning Annie (5.685728); The Romantics (film) (5.416189)

[Component 3/3] Query Processor (a) VSM ltc.ltc (4pts)
[Task 7/12] Testing VSM ltc.ltc...
5 results, top score: 0.195387 | Top 3: I Love You, I Love You Not (0.195387); Crash Pad (0.190719); Boy Culture (0.178577)

[Component 3/3] Query Processor (b) SMART Variations (+2pts)
[Task 8/12] Testing VSM etc.ltc...
5 results, top score: 0.195387 | Top 3: I Love You, I Love You Not (0.195387); Crash Pad (0.190719); Boy Culture (0.178577)
[Task 9/12] Testing VSM inc.ltc...
5 results, top score: 0.231269 | Top 3: Champions: A Love Story (0.231269); I Love You, I Love You Not (0.192966); Boy Culture (0.176179)
[Task 10/12] Testing VSM atc.ltc...
5 results, top score: 0.195387 | Top 3: I Love You, I Love You Not (0.195387); Crash Pad (0.190719); Boy Culture (0.178577)

[Component 3/3] Query Processor (c) BM25/LM Extension (+2pts total)
[Task 11/12] Testing BM25 (+1pt)...
10 results, top score: 11.415098 | Top 3: I Love You, I Love You Not (11.415098); Peter and Vandy (10.851616); Boy Culture (10.251861)

[Task 12/12] Testing Language Model (+1pt)...
5 results, top score: -10.005406 | Top 3: Isn't It Romantic (2019 film) (-10.005406); The Lost Valentine (-10.379781); God of Love (film) (-10.381941)
```

Figure 2: Dynamic index operations demonstrating document insertion, update, and deletion with auxiliary index and tombstone management.

A.3 Complete System Timeline

=====			
TIMELINE SUMMARY - Assignment Components			
=====			
1. Tokenizer (3pts)			
Task	Operation	Query	Output
✓ 1. Tokenizer (3pts)	tokenize	A space adventure about alie...	7 tokens, 7 unique: space adventur alien explor ...
2. Indexer (a) Memory (5pts)			
Task	Operation	Query	Output
✓ 2. Indexer RAM SPIMI (5pts)	build --mode=memory	N/A	Indexed 17838 documents, 65429 terms, 2985645 to...
2. Indexer (b) Disk Extension (+2pts)			
Task	Operation	Query	Output
✓ 3. Indexer Disk (+2pts)	build --mode=disk	N/A	Indexed 17838 documents, 65429 terms in lexicon,...
2. Indexer (c) Updatable Extension (+2pts)			
Task	Operation	Query	Output
✓ 4. Updatable Insert (+1pt)	add --mode=updatable	space exploration	Added doc ID 17830: 'Test Film' Plot: "A tes...
✓ 5. Updatable Update (+0.5pt)	update --mode=updatable	time travel	Updated doc 11262: 'Timequest (film)' → new ID 1...
✓ 6. Updatable Delete (+0.5pt)	delete --mode=updatable	romantic	Deleted doc ID 100 (in deleted set: True) Sear...
3. Query Processor (a) VSM ltc.ltc (4pts)			
Task	Operation	Query	Output
✓ 7. VSM ltc.ltc (4pts)	search --test=ltc	romantic love story	5 results, top score: 0.195387 Top 3: I Love Y...
3. Query Processor (b) SMART Variations (+2pts)			
Task	Operation	Query	Output
✓ 8. VSM ntc.ltc (+0.67pt)	search --test=ntc	romantic love story	5 results, top score: 0.195387 Top 3: I Love Y...
✓ 9. VSM lnc.ltc (+0.67pt)	search --test=lnc	romantic love story	5 results, top score: 0.231269 Top 3: Champion...
✓ 10. VSM atc.ltc (+0.66pt)	search --test=atc	romantic love story	5 results, top score: 0.195387 Top 3: I Love Y...
3. Query Processor (c) BM25/LM Extension (+2pts)			
Task	Operation	Query	Output
✓ 11. BM25 (+1pt)	search --test=bm25	romantic love story	10 results, top score: 11.415898 Top 3: I Love...
✓ 12. Language Model (+1pt)	search --test=lm	romantic love story	5 results, top score: -18.055406 Top 3: Isn't ...
=====			

Figure 3: Final timeline verification showing all 12 assignment tasks completed successfully, including tokenization (3pts), SPIMI indexing (5pts), disk index (+2pts), updatable index (+2pts), VSM ltc.ltc (4pts), SMART variations (+2pts), and BM25 (+2pts).

Summary: The three figures above provide visual evidence of complete system functionality. Figure 1 demonstrates correct index construction across all three modes (memory, disk, updatable) with consistent vocabulary statistics. Figure 2 verifies dynamic operations including insert, update, and delete with proper ranking maintenance. Figure 3 confirms that all assignment requirements were met, with 12 out of 12 tasks passing successfully within the Docker environment.